Inclass - Lab (Week 3)



About the data set (Life Expectancy data)

The dataset contains information about several health and economic factors that contribute to average life expectancy for different countries. Attribute information:

Country: Name of the country

Year: Year the observations were recorded

Status: Whether the country is Developed or Developing

Adult_Mortality: Mortality rate for age group 15-60 out of every 1000 individuals of the population

Infant_Deaths: Number of infant deaths per 1000 population

Alcohol: Alcohol, recorded per capita (15+) consumption (in litres of pure alcohol)

Hepatitis B: Hepatitis B (HepB) immunization coverage for 1 year olds (Percentage)

Measles: Number of reported cases for measles per 1000 from population

BMI: Average Body Mass Index for entire population

Underfive_Deaths: Number of deaths under 5 years of age per 1000 population

Polio: Polio (Pol3) immunization coverage for 1 year olds (Percentage)

Diphtheria: Diphtheria tetanus toxoid and pertussis (DTP3) immunization coverage for 1 year olds (Percentage)

HIV: Deaths per 1000 live births due to HIV/AIDS (0-4 years)

GDP: Gross Domestic Product per capita (in USD)

Population: Population of the country

Malnourished10_19: Prevalence of malnutrition among children and adolescents for Age 10 to 19 (Percentage)

Malnourished5_9: Prevalence of malnutrition among children for Age 5 to 9 (Percentage)

Income_Index: Human Development Index (HDI) in terms of national income per capita (index ranging from 0 to 1)

Schooling: Number of years of Schooling

Life_Expectancy: Life Expectancy in age for the country

Table of Content

- 1. Data Preparation
- 2. Linear Regression
- 3. Gradient Descent
- 4. Regularization
- 5. Grid Search

Import the required libraries

In []: |# type your code here

```
In [1]: # import 'Pandas'
        import pandas as pd
        # import 'Numpy'
        import numpy as np
        # import subpackage of Matplotlib
        import matplotlib.pyplot as plt
        # import 'Seaborn'
        import seaborn as sns
        # to suppress warnings
        from warnings import filterwarnings
        filterwarnings('ignore')
        # display all columns of the dataframe
        pd.options.display.max_columns = None
        # display all rows of the dataframe
        pd.options.display.max_rows = None
        # to display the float values upto 6 decimal places
        pd.options.display.float_format = '{:.6f}'.format
        # import train-test split
        from sklearn.model_selection import train_test_split
        # import various functions from statsmodels
        import statsmodels
        import statsmodels.api as sm
        # import 'stats'
        from scipy import stats
        # 'metrics' from sklearn is used for evaluating the model performance
        from sklearn.metrics import mean_squared_error
        from sklearn.metrics import mean_absolute_error
        # import function to perform linear regression
        from sklearn.linear_model import LinearRegression
        # import StandardScaler to perform scaling
        from sklearn.preprocessing import StandardScaler
        # import SGDRegressor from sklearn to perform linear regression with stochastic gradient descent
        from sklearn.linear_model import SGDRegressor
        # import function for ridge regression
        from sklearn.linear_model import Ridge
        # import function for lasso regression
        from sklearn.linear_model import Lasso
        # import function for elastic net regression
        from sklearn.linear_model import ElasticNet
        # import function to perform GridSearchCV
        from sklearn.model_selection import GridSearchCV
```

```
In [2]: # set the plot size using 'rcParams'
# once the plot size is set using 'rcParams', it sets the size of all the forthcoming plots in the file
# pass width and height in inches to 'figure.figsize'
plt.rcParams['figure.figsize'] = [15,8]
```

Load the Life expectancy dataset and display the first five records

```
In [16]: # load the csv file
          # type your code here
          df = pd.read_csv('Life_Expectancy.csv')
          # display the first five observations
          # type your code here
          df.head()
Out[16]:
                                                                            Hepatitis
                Country Year
                                 Status Adult_Mortality Infant_Deaths
                                                                   Alcohol
                                                                                     Measles
                                                                                                  BMI Underfive_Deaths Polio Diphtheria
                                                                                  В
           0 Afghanistan 2015 Developing
                                           263.000000
                                                               62 0.010000 65.000000
                                                                                        1154 19.100000
                                                                                                                                   65 0.10
                                                                                                                    83
                                                                                                                          6
                                            74.000000
           1
                 Albania 2015 Developing
                                                                0 4.600000 99.000000
                                                                                           0 58.000000
                                                                                                                    0
                                                                                                                         99
                                                                                                                                   99 0.10
           2
                                            19.000000
                                                               21
                                                                                          63 59.500000
                                                                                                                         95
                 Algeria 2015 Developing
                                                                      NaN 95.000000
                                                                                                                    24
                                                                                                                                   95 0.10
                                                                                                                                   64 1.90
                 Angola 2015 Developing
                                           335.000000
                                                               66
                                                                      NaN 64.000000
                                                                                         118 23.300000
                                                                                                                    98
                                                                                                                          7
                                            13.000000
                                                                0
                                                                      NaN 99.000000
                                                                                           0 47.700000
                                                                                                                    0
                                                                                                                         86
                                                                                                                                   99 0.20
                Antigua 2015 Developing
          Check the shape of the data and display its information
 In [4]: # check the size of the dataframe
          df.shape
 Out[4]: (182, 20)
 In [5]: | # display dataset information understand the dataset
          # type your code here
          df.info()
          <class 'pandas.core.frame.DataFrame'>
          RangeIndex: 182 entries, 0 to 181
          Data columns (total 20 columns):
                                    Non-Null Count Dtype
           #
               Column
          _ _ _
               -----
                                    -----
                                    182 non-null
           0
               Country
                                                     object
           1
               Year
                                    182 non-null
                                                     int64
                                    182 non-null
           2
               Status
                                                     object
           3
               Adult_Mortality
                                    173 non-null
                                                     float64
           4
               Infant_Deaths
                                    182 non-null
                                                     int64
               Alcohol
                                    15 non-null
                                                     float64
           5
           6
               Hepatitis B
                                    173 non-null
                                                     float64
           7
               Measles
                                    182 non-null
                                                     int64
           8
               BMI
                                    178 non-null
                                                     float64
           9
                                   182 non-null
                                                     int64
               Underfive_Deaths
                                    182 non-null
                                                     int64
               Polio
           10
               Diphtheria
                                    182 non-null
                                                     int64
           11
           12
               HIV
                                    182 non-null
                                                     float64
```

Let's begin with some hands-on practice exercises

157 non-null

141 non-null

178 non-null

169 non-null

173 non-null

173 non-null

Malnourished10_19 178 non-null

dtypes: float64(12), int64(6), object(2)

float64

float64

float64

float64

float64

float64

float64

1. Data Preparation



GDP

Population

Malnourished5_9

Life_Expectancy

Income Index

memory usage: 28.6+ KB

Schooling

13

14

15

17

18

19

1. The dataset gives the life expectancy for different countries. Check if any information about life expectancy is missing from the given records. How do we fix the problem if present?

```
In [27]: # type your code here
df['Life_Expectancy'].isnull().sum()
Out[27]: 0
In [18]: na_data = df['Life_Expectancy'].isnull()
```

```
In [20]: | na_data.index[na_data == True].tolist()
Out[20]: [46, 100, 104, 111, 118, 122, 135, 138, 168]
In [25]: df = df.drop(df.index[[46, 100, 104, 111, 118, 122, 135, 138, 168]])
In [28]: df.shape
Out[28]: (173, 20)
           2. How do we handle the missing values from the entire dataset?
In [30]: # type your code here
          df.isnull().sum()
Out[30]: Country
                                   0
                                   0
          Year
          Status
                                   0
          Adult_Mortality
                                   0
          Infant_Deaths
                                   0
          Alcohol
                                166
          Hepatitis B
                                   9
          Measles
          BMI
                                   2
          Underfive_Deaths
          Polio
                                   0
          Diphtheria
          HIV
                                   0
          GDP
                                  21
                                  34
          Population
          Malnourished10_19
                                  2
          Malnourished5_9
                                   2
          Income_Index
                                   7
                                   7
          Schooling
          Life_Expectancy
          dtype: int64
In [46]: df_missing_values = df[['Hepatitis B','BMI','GDP','Population','Malnourished10_19',
                                            'Malnourished5_9','Income_Index','Schooling']]
          # plot histogram of all variables which have missing values
          # set the number of bins to 20
          # set the figure size using 'figsize'
          df_missing_values.hist(bins = 20, figsize = (15,8))
          # adjust the subplots
          plt.tight_layout()
          # display the plot
          plt.show()
                            Hepatitis B
                                                                                                                   GDP
                                                                         BMI
                                                                                               100
                                                     25
           60
                                                                                               80
                                                     20
                                                     15
                                                                                               60
                                                                                               40
                                                     10
           20
                                                                                               20
                                                                                                      10000 20000 30000 40000 50000 60000
                                                                                       70
                                                            10
                                                                20
                                                                     30
                                                                         40
                                                                                           80
                                                                                                              Malnourished5 9
                             Population
                                                                   Malnourished10_19
           150
                                                     40
                                                     30
           100
           50
                                              2.5
1e8
                                                                       Schooling
                            Income_Index
           15
                                                     20
           10
                                                                   10
                                        0.8
                                                                                 16
                       0.5
                             0.6
                                  0.7
                                              0.9
```

```
In [34]: | df['Hepatitis B'] = df['Hepatitis B'].fillna(df['Hepatitis B'].median())
In [35]: |df['BMI'] = df['BMI'].fillna(df['BMI'].mean())
In [36]: df['GDP'] = df['GDP'].fillna(df['GDP'].median())
In [37]: |df['Population'] = df['Population'].fillna(df['Population'].median())
In [38]: | df['Malnourished10_19'] = df['Malnourished10_19'].fillna(df['Malnourished10_19'].median())
In [39]: df['Malnourished5_9'] = df['Malnourished5_9'].fillna(df['Malnourished5_9'].median())
In [40]: | df['Income_Index'] = df['Income_Index'].fillna(df['Income_Index'].mean())
In [41]: | df['Schooling'] = df['Schooling'].fillna(df['Schooling'].mean())
In [42]: | df.isnull().sum()
Out[42]: Country
                                 0
                                 0
         Year
                                 0
         Status
         Adult_Mortality
         Infant_Deaths
         Alcohol
                               166
         Hepatitis B
                                 0
         Measles
                                 0
         BMI
         Underfive_Deaths
                                 0
         Polio
         Diphtheria
         HIV
         GDP
                                 0
                                 0
         Population
         Malnourished10_19
                                 0
         Malnourished5_9
                                 0
                                 0
         Income_Index
                                 0
         Schooling
         Life_Expectancy
         dtype: int64
In [44]: | df.drop('Alcohol', axis=1, inplace=True)
In [45]: | df.isnull().sum()
Out[45]: Country
                               0
         Year
                               0
                               0
         Status
                               0
         Adult_Mortality
         Infant_Deaths
         Hepatitis B
                               0
         Measles
                               0
         BMI
                               0
         Underfive_Deaths
         Polio
                               0
         Diphtheria
                               0
         HIV
                               0
         GDP
                               0
         Population
                               0
         Malnourished10_19
                               0
         Malnourished5_9
         Income_Index
         Schooling
                               0
         Life_Expectancy
         dtype: int64
          3. Are there any redundant features in the data?
In [47]: # type your code here
         df.describe(include='object')
Out[47]:
                   Country
                              Status
                       173
                                173
           count
```

unique

freq

173

top Afghanistan Developing

2

141

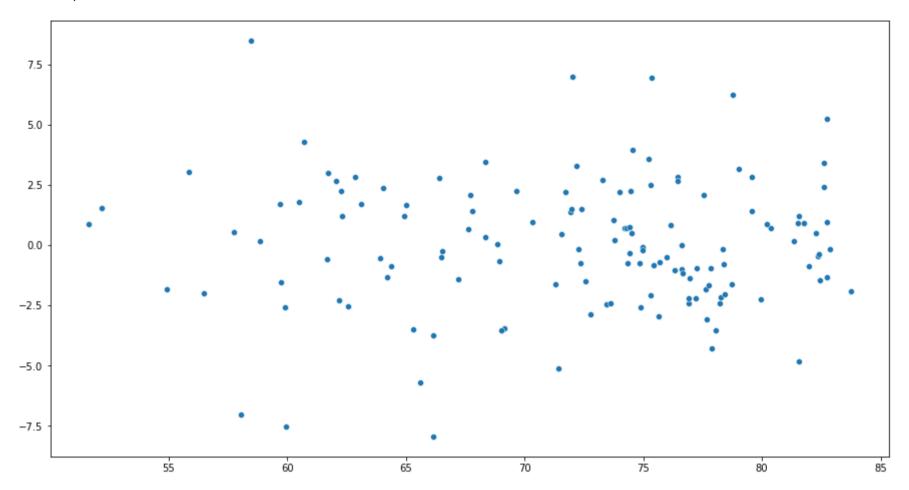
```
In [48]: | df.drop('Country', axis=1, inplace=True)
          df.describe()
In [49]:
Out[49]:
                                                                                                                             Diphtheria
                         Year Adult_Mortality Infant_Deaths Hepatitis B
                                                                            Measles
                                                                                                Underfive_Deaths
                                                                                                                      Polio
                                                                                                                                               HIV
                                                                         173.000000
                                                                                                                 173.000000
                   173.000000
                                   173.000000
                                                173.000000
                                                           173.000000
                                                                                    173.000000
                                                                                                      173.000000
                                                                                                                             173.000000
                                                                                                                                        173.000000
            count
                                   149.971098
                                                                                                                             84.508671
                  2014.988439
                                                 22.872832
                                                             82.716763
                                                                        1559.208092
                                                                                     43.466667
                                                                                                       30.352601
                                                                                                                  83.040462
                                                                                                                                          0.647399
            mean
                     0.152057
                                   95.679846
                                                 84.268577
                                                             24.066903
                                                                        8055.443243
                                                                                                                             22.687364
              std
                                                                                     20.378109
                                                                                                      108.428722
                                                                                                                  24.833951
                                                                                                                                          1.361795
                  2013.000000
                                                             6.000000
                                                                           0.000000
                                                                                                        0.000000
                                    1.000000
                                                  0.000000
                                                                                      2.500000
                                                                                                                   5.000000
                                                                                                                              6.000000
                                                                                                                                          0.100000
             min
                  2015.000000
                                   74.000000
                                                  0.000000
                                                             81.000000
                                                                           0.000000
                                                                                                                  84.000000
                                                                                                                             84.000000
                                                                                                                                          0.100000
             25%
                                                                                     24.400000
                                                                                                        0.000000
             50%
                  2015.000000
                                   137.000000
                                                  2.000000
                                                             93.000000
                                                                          16.000000
                                                                                     51.000000
                                                                                                        3.000000
                                                                                                                  93.000000
                                                                                                                             94.000000
                                                                                                                                          0.100000
                  2015.000000
                                   199.000000
                                                                                                                                          0.300000
             75%
                                                 17.000000
                                                             97.000000
                                                                         212.000000
                                                                                     61.600000
                                                                                                       21.000000
                                                                                                                  97.000000
                                                                                                                             97.000000
             max 2015.000000
                                                                       90387.000000
                                  484.000000
                                                                                                     1100.000000
                                                                                                                                          9.300000
                                                910.000000
                                                             99.000000
                                                                                     77.600000
                                                                                                                  99.000000
                                                                                                                             99.000000
In [50]: |df['Year'].value_counts()
Out[50]: 2015
                    172
           2013
                      1
           Name: Year, dtype: int64
In [51]: |df.drop('Year', axis=1, inplace=True)
In [52]: | df.drop('Infant_Deaths', axis=1, inplace=True)
In [53]:
          df.shape
Out[53]: (173, 16)
                 4. Perform dummy encoding for appropriate variables of the dataset if required
In [54]: # type your code here
           df['Status'] = pd.get_dummies(data=df['Status'],drop_first=True)
In [55]:
          df.head()
Out[55]:
                                      Hepatitis
              Status Adult_Mortality
                                                                                                         HIV
                                                                                                                      GDP
                                                             ВМІ
                                                                  Underfive_Deaths Polio Diphtheria
                                               Measles
                                                                                                                                 Population Malnouri
           0
                          263.000000
                                     65.000000
                                                  1154
                                                        19.100000
                                                                                83
                                                                                       6
                                                                                                    0.100000
                                                                                                                584.259210
                                                                                                                           33736494.000000
                          74.000000 99.000000
                                                     0 58.000000
                                                                                 0
                                                                                                 99 0.100000
                                                                                                               3954.227830
                                                                                      99
                                                                                                                               28873.000000
           1
                           19.000000 95.000000
                                                    63 59.500000
                                                                                24
                                                                                      95
                                                                                                 95
                                                                                                    0.100000
                                                                                                               4132.762920
                                                                                                                           39871528.000000
                                                    118 23.300000
                                                                                                               3695.793748
                          335.000000 64.000000
                                                                                98
                                                                                       7
                                                                                                    1.900000
                                                                                                                             2785935.000000
                   1
                                                                                                 64
                           13.000000 99.000000
                                                     0 47.700000
                                                                                 0
                                                                                      86
                                                                                                 99 0.200000
                                                                                                              13566.954100
                                                                                                                             2174645.000000
           2. Linear Regression
                   5. Build a full model on the given data. Check whether the obtained residuals have constant variance
In [56]:
          # type your code here
          X = df.iloc[:,:15]
          X = sm.add\_constant(X)
          y = df['Life_Expectancy']
```

X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=10, test_size=0.2)

MLR_model = sm.OLS(y_train, X_train).fit()

```
In [58]: y_pred = MLR_model.fittedvalues
    residuals = MLR_model.resid
    sns.scatterplot(x=y_pred,y= residuals)
```

Out[58]: <AxesSubplot:>



In []: # There is no pattern seen in the scatterplot hence the data contains constant variance.

6. Build a model to study the impact of immunization on life expectancy, using least squares regression. Find the unexplained variation for the model

```
In [59]: # type your code here
X = df.iloc[:,[2,6,7]]
y = df['Life_Expectancy']

linreg = LinearRegression()
MLR_model = linreg.fit(X,y)
y_pred = MLR_model.predict(X)
residual = np.array(y - y_pred)

unexplained_var = np.sum(residual**2)
print(unexplained_var)
```

7272.4649163342965



7. Build a model to study the impact of malnutrition on life expectancy, using least squares regression. Interpret the coefficients

```
In [60]: # type your code here
          X = df.iloc[:,[11,12]]
          X = sm.add\_constant(X)
          y = df['Life_Expectancy']
          X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=10, test_size=0.2)
          MLR_model = sm.OLS(y_train, X_train).fit()
          MLR_model.summary()
Out[60]:
           OLS Regression Results
               Dep. Variable:
                              Life_Expectancy
                                                   R-squared:
                                                                 0.226
                      Model:
                                        OLS
                                               Adj. R-squared:
                                                                 0.214
                    Method:
                                Least Squares
                                                   F-statistic:
                                                                 19.66
                       Date: Sun, 19 Mar 2023 Prob (F-statistic): 3.20e-08
                                               Log-Likelihood:
                      Time:
                                     11:15:31
                                                                -466.03
            No. Observations:
                                         138
                                                         AIC:
                                                                 938.1
                Df Residuals:
                                         135
                                                         BIC:
                                                                 946.8
                                           2
                   Df Model:
            Covariance Type:
                                   nonrobust
                                 coef std err
                                                      P>|t|
                                                            [0.025 0.975]
                        const 76.3070
                                        0.916 83.320 0.000
                                                           74.496 78.118
            Malnourished10_19
                                              -0.654 0.514
                              -0.5024
                                        0.769
                                                            -2.023
                                                                    1.018
              Malnourished5_9 -0.3930
                                        0.754
                                              -0.521 0.603
                                                            -1.884
                                                                    1.098
                 Omnibus: 6.049
                                    Durbin-Watson:
                                                    1.840
            Prob(Omnibus): 0.049
                                  Jarque-Bera (JB):
                                                    6.225
                    Skew: -0.512
                                         Prob(JB): 0.0445
                  Kurtosis: 2.820
                                         Cond. No.
                                                      16.1
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

3. Gradient Descent

8. Build a model to study the influence of adult mortality rate on life expectancy using both the least squares regression and gradient descent approach. Are the fits for both the models different?

```
In [61]: # type your code here
    X = df.iloc[:,1].values
    X = X.reshape(-1,1)

# set the dependent variable
    y = df['Life_Expectancy']

# initialize the standard scalar
    X_scaler = StandardScaler()

# standardize all the columns of df_LifeExp
    X = X_scaler.fit_transform(X)

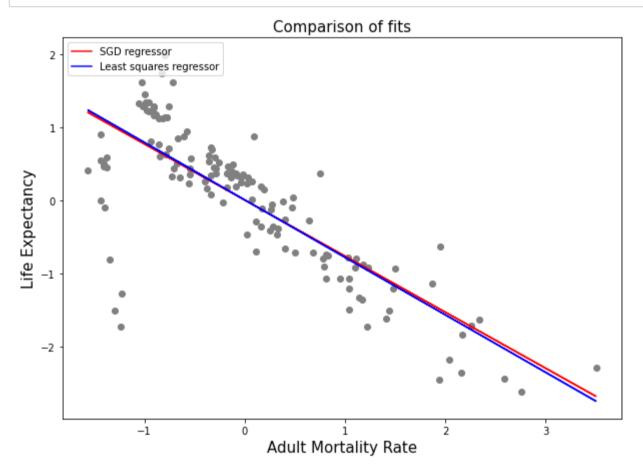
# standardize the target explicitly and storing it in a new variable 'y'
    y = (y - y.mean()) / y.std()

# split data into train subset and test subset
    # set 'random_state' to generate the same dataset each time you run the code
    # 'test_size' returns the proportion of data to be included in the test set
    X_train, X_test, y_train, y_test = train_test_split(X, y, random_state = 10, test_size = 0.2)
```

```
In [62]: linreg = LinearRegression()
MLR_model = linreg.fit(X_train,y_train)
linpred = MLR_model.predict(X_train)
```

```
In [63]: sgd = SGDRegressor()
MLR_model_sgd = sgd.fit(X_train,y_train)
sgdpred = MLR_model_sgd.predict(X_train)
```

```
In [64]: # set the figure size
         plt.figure(figsize=(10,7))
         # plot the scatter plot
         # colors: set the color of the points in the scatter plot
         plt.scatter(X_train, y_train,color='grey')
         # set xlabel for the plot
         plt.xlabel("Adult Mortality Rate",fontsize = 15)
         # set ylabel for the plot
         plt.ylabel("Life Expectancy", fontsize = 15)
         # set title for the plot
         plt.title("Comparison of fits", fontsize = 15)
         # plot the regression line for the two models
         # color: set the color of the line
         # label: set the label of the line for the legend
         plt.plot(X_train,sgdpred,color='red',label='SGD regressor')
         plt.plot(X_train,linpred,color='blue',label='Least squares regressor')
         # set the position of legend
         plt.legend(loc='upper left')
         # disply the plot
         plt.show()
```



9. Build a model to analyze the influence of immunization on life expectancy. Use SGD and plot a horizontal multiple barchart to compare the values of beta coefficients with values obtained by the full model build least squares regression.

```
In [68]: # type your code here
X = df.iloc[:,[2,6,7]]

# set the dependent variable
y = df['Life_Expectancy']

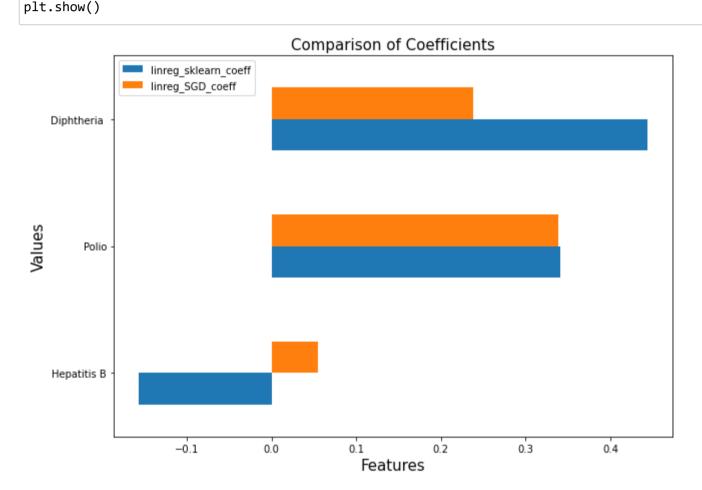
# initialize the standard scalar
X_scaler = StandardScaler()

# standardize all the columns of df_lifeExp
X = X_scaler.fit_transform(X)

# standardize the target explicitly and storing it in a new variable 'y'
y = (y - y.mean()) / y.std()

# split data into train subset and test subset
# set 'random_state' to generate the same dataset each time you run the code
# 'test_size' returns the proportion of data to be included in the test set
X_train, X_test, y_train, y_test = train_test_split(X, y, random_state = 10, test_size = 0.2)
```

```
In [69]: sgd1 = SGDRegressor()
MLR_model_sgd1 = sgd1.fit(X_train,y_train)
```



set the x-axis labels horizontally
plt.xticks(rotation = 'horizontal')

set the position of Legend
fig.legend(loc='upper left')

display the plot

In []: # For polio both the model interpretations are nearly similar
For diptheria the linreg coeff are higher than SGD.
For hepatitis B the SGD coeff are higher and the lin reg coeff have gone towards negative.

10. Build a full model using least squares regression. Check whether the model overfits the training data or not.

```
In [78]: # type your code here
X = df.iloc[:,:15]
y = df['Life_Expectancy']

X_train, X_test, y_train, y_test = train_test_split(X, y, random_state = 10, test_size = 0.2)

linreg = LinearRegression()
MLR_model2 = linreg.fit(X_train,y_train)
```

```
In [79]: # Train
train_pred = MLR_model2.predict(X_train)

mse_train = mean_squared_error(y_train,train_pred)
rmse_train = round(np.sqrt(mse_train),4)
```

```
In [80]: # Test
    test_pred = MLR_model2.predict(X_test)

mse_test = mean_squared_error(y_test,test_pred)
    rmse_test = round(np.sqrt(mse_test),4)
```

```
print(mse_test,rmse_test)
              7.0590856828921575 2.6569
             11.72600488499266 3.4243
              11. Can we use a linear regression model to analyze how all features from the dataset impact life expectancy?
 In [ ]: # type your code here
              # Yes we can use linear regression because the target variable is numerical and is dependent on independent variables
In [85]: | df_features = df.iloc[:,:15]
              sns.heatmap(df_features.corr(),annot=True)
Out[85]: <AxesSubplot:>
                                                                                                                                                                               - 1.0
                            Status
                                                      -0.16
                                                              0.085
                                                                       -0.29
                                                                               0.13
                                                                                        -0.23
                                                                                                -0.23
                                                                                                         0.19
                                                                                                                 -0.34
                                                                                                                        0.0048
                                                                                                                                           0.38
                                                                                                                                                   -0.54
                                                                                                                                                            -0.53
                                      1
                   Adult_Mortality
                                               1
                                                      -0.25
                                                              0.055
                                                                       -0.4
                                                                                        -0.39
                                                                                                -0.33
                                                                                                                 -0.32
                                                                                                                         0.059
                                                                                                                                            0.3
                                                                                                                                                   -0.65
                                                                                                                                                            -0.54
                                                                                                                                                                               - 0.8
                                              -0.25
                                                       1
                                                                               -0.11
                                                                                                0.92
                                                                                                         -0.4
                       Hepatitis B
                                     -0.16
                                                              0.014
                                                                                                                  0.13
                                                                                                                         -0.069
                                                                                                                                  -0.079
                                                                                                                                           -0.12
                                                                                                                                                            0.34
                                             0.055
                                                      0.014
                                                               1
                                                                       -0.14
                                                                                0.78
                                                                                        -0.026
                                                                                                        -0.027
                                                                                                                 -0.075
                         Measles
                                                                                               4.1e-05
                                                                                                                          0.13
                                                                                                                                           0.34
                                                                                                                                                   -0.081
                                                                                                                                                           -0.073
                                                                                                                                                                               0.6
                              BMI
                                     -0.29
                                              -0.4
                                                       0.2
                                                               -0.14
                                                                        1
                                                                               -0.23
                                                                                        0.23
                                                                                                         -0.29
                                                                                                                 0.31
                                                                                                                         -0.026
                                                                                                                                  -0.47
                                                                                                                                           -0.49
                                                      -0.11
                                                               0.78
                                                                       -0.23
                                                                                 1
                                                                                                                                                   -0.24
                                                                                                                                                                               - 0.4
                Underfive_Deaths
                                     0.13
                                              0.2
                                                                                        -0.15
                                                                                                -0.14
                                                                                                         0.12
                                                                                                                 -0.12
                                                                                                                                                            -0.24
                             Polio
                                     -0.23
                                              -0.39
                                                              -0.026
                                                                       0.23
                                                                               -0.15
                                                                                         1
                                                                                                0.66
                                                                                                         -0.41
                                                                                                                  0.22
                                                                                                                          -0.23
                                                                                                                                   -0.2
                                                                                                                                            -0.2
                                                                                                                                                                                0.2
                                              -0.33
                                                      0.92
                                                             4.1e-05
                                                                       0.22
                                                                                                  1
                                                                                                         -0.45
                       Diphtheria
                                     -0.23
                                                                               -0.14
                                                                                        0.66
                                                                                                                  0.21
                                                                                                                         -0.073
                                                                                                                                  -0.12
                                                                                                                                           -0.16
                                                      -0.4
                                                              -0.027
                                                                       -0.29
                                                                                        -0.41
                                                                                                -0.45
                                                                                                          1
                                                                                                                 -0.18
                                                                                                                         0.044
                                                                                                                                           0.15
                                                                                                                                                   -0.48
                                                                                                                                                            -0.39
                              HIV
                                     0.19
                                                                               0.12
                                                                                                                                                                               0.0
                                     -0.34
                                              -0.32
                                                      0.13
                                                                       0.31
                                                                               -0.12
                                                                                        0.22
                                                                                                0.21
                                                                                                         -0.18
                                                                                                                   1
                                                                                                                          0.033
                                                                                                                                  -0.23
                                                                                                                                           -0.23
                             GDP
                                                              -0.075
                        Population
                                    0.0048
                                             0.059
                                                      -0.069
                                                               0.13
                                                                      -0.026
                                                                                0.32
                                                                                        -0.23
                                                                                                -0.073
                                                                                                        0.044
                                                                                                                 0.033
                                                                                                                           1
                                                                                                                                  0.0091
                                                                                                                                         0.0034
                                                                                                                                                   -0.02
                                                                                                                                                           0.0003
                                                                                                                                                                               -0.2
               Malnourished10 19
                                              0.29
                                                      -0.079
                                                               0.35
                                                                       -0.47
                                                                                        -0.2
                                                                                                -0.12
                                                                                                         0.17
                                                                                                                 -0.23
                                                                                                                         0.0091
                                                                                                                                   1
                                                                                                                                           0.97
                                                                                                                                                    -0.5
                                                                                                                                                            -0.48
                                                                                                                 -0.23
                 Malnourished5_9
                                     0.38
                                                      -0.12
                                                               0.34
                                                                       -0.49
                                                                                        -0.2
                                                                                                -0.16
                                                                                                         0.15
                                                                                                                         0.0034
                                                                                                                                  0.97
                                                                                                                                            1
                                                                                                                                                    -0.5
                                                                                                                                                            -0.48
                                                                                                                                                                               - -0.4
                                     -0.54
                                                                                                                                                     1
                    Income_Index
                                              -0.65
                                                      0.34
                                                              -0.081
                                                                               -0.24
                                                                                                         -0.48
                                                                                                                          -0.02
                                                                                                                                   -0.5
                                                                                                                                            -0.5
                                                                                                                                                            0.92
                        Schooling
                                     -0.53
                                              -0.54
                                                              -0.073
                                                                                -0.24
                                                                                                         -0.39
                                                                                                                         0.0003
                                                                                                                                  -0.48
                                                                                                                                           -0.48
                                                                                         Polio
                                      Status
                                                                                                                                   Malnourished10_19
                                                                                                                                            Malnourished5 9
                                                                        ₩
                                                                                                          ⋛
                                                                                                                           Population
                                                                                                                   ලි
                                              Adult Mortality
                                                                                 Underfive_Deaths
                                                                                                 Diphtheria
                                                                                                                                                    Income_Index
                                                       Hepatitis
```

4. Regularization

In [81]: |print(mse_train,rmse_train)

12. Can we perform regression analysis without eliminating features involved in multicollinearity detected in question 11?

```
In [90]: # type your code here
X = df.iloc[:,:15]

# set the dependent variable
y = df['Life_Expectancy']

# initialize the standard scalar
X_scaler = StandardScaler()

# standardize all the columns of df_lifeExp
X = X_scaler.fit_transform(X)

# standardize the target explicitly and storing it in a new variable 'y'
y = (y - y.mean()) / y.std()

# split data into train subset and test subset
# set 'random_state' to generate the same dataset each time you run the code
# 'test_size' returns the proportion of data to be included in the test set
X_train, X_test, y_train, y_test = train_test_split(X, y, random_state = 110, test_size = 0.2)
```

```
In [91]: ridge = Ridge(alpha=5.5)
         ridge_model = ridge.fit(X_train, y_train)
In [92]: # Train
         train_pred = ridge_model.predict(X_train)
         mse_train = mean_squared_error(y_train,train_pred)
         rmse_train = round(np.sqrt(mse_train),4)
         print(rmse_train)
         0.3528
In [93]: # Test
         test_pred = ridge_model.predict(X_test)
         mse_test = mean_squared_error(y_test,test_pred)
         rmse_test = round(np.sqrt(mse_test),4)
         print(rmse_test)
         0.3533
 In [ ]: # Ridge
                 13. Build a model to study the impact of diseases and immunization on life expectancy. Identify the significant variables
                 from the model. Use 'alpha = 1.5' if required.
In [94]: # type your code here
         # type your code here
         X = df.iloc[:,[2,3,6,7,8]]
         # set the dependent variable
         y = df['Life_Expectancy']
         # initialize the standard scalar
         X_scaler = StandardScaler()
         # standardize all the columns of df_lifeExp
         X = X_scaler.fit_transform(X)
         # standardize the target explicitly and storing it in a new variable 'y'
         y = (y - y.mean()) / y.std()
         # split data into train subset and test subset
         # set 'random_state' to generate the same dataset each time you run the code
         # 'test_size' returns the proportion of data to be included in the test set
         X_train, X_test, y_train, y_test = train_test_split(X, y, random_state = 10, test_size = 0.2)
In [95]: ridge = Ridge(alpha=1.5)
         ridge_model = ridge.fit(X_train, y_train)
In [96]: # Train
         train_pred = ridge_model.predict(X_train)
         mse_train = mean_squared_error(y_train,train_pred)
         rmse_train = round(np.sqrt(mse_train),4)
         print(rmse_train)
         0.7101
In [97]: # Test
         test_pred = ridge_model.predict(X_test)
         mse_test = mean_squared_error(y_test,test_pred)
         rmse_test = round(np.sqrt(mse_test),4)
         print(rmse_test)
         0.7136
In [98]: X = df.iloc[:,[2,3,6,7,8]]
         X = sm.add_constant(X)
         y = df['Life_Expectancy']
         MLR_model3 = sm.OLS(y_train,X_train).fit()
```

```
In [100]: # Train
           train_pred = MLR_model3.predict(X_train)
           mse_train = mean_squared_error(y_train,train_pred)
           rmse_train = round(np.sqrt(mse_train),4)
           print(rmse_train)
           0.7102
In [101]:
          # Test
           test_pred = MLR_model3.predict(X_test)
           mse_test = mean_squared_error(y_test,test_pred)
           rmse_test = round(np.sqrt(mse_test),4)
           print(rmse_test)
           0.7117
In [102]:
          MLR_model3.summary()
Out[102]:
           OLS Regression Results
                Dep. Variable:
                              Life_Expectancy
                                                 R-squared (uncentered):
                                                                         0.498
                      Model:
                                       OLS Adj. R-squared (uncentered):
                                                                         0.479
                     Method:
                                Least Squares
                                                            F-statistic:
                                                                         26.39
                       Date: Sun, 19 Mar 2023
                                                      Prob (F-statistic): 1.87e-18
                      Time:
                                    12:18:25
                                                        Log-Likelihood:
                                                                       -148.59
            No. Observations:
                                                                 AIC:
                                        138
                                                                         307.2
                                                                 BIC:
                Df Residuals:
                                        133
                                                                         321.8
                   Df Model:
                                          5
             Covariance Type:
                                   nonrobust
                  coef std err
                                   t P>|t| [0.025 0.975]
            x1 -0.1130
                        0.138 -0.821 0.413 -0.385
                                                  0.159
                -0.0663
                        0.057 -1.168 0.245 -0.179
                                                   0.046
                0.2707
                        0.081
                               3.355 0.001
                                            0.111
                                                   0.430
            x3
                0.2251
                        0.065 -6.673 0.000 -0.563 -0.306
            х5
               -0.4345
                 Omnibus: 7.053
                                   Durbin-Watson:
                                                   2.005
            Prob(Omnibus): 0.029 Jarque-Bera (JB):
                                                   6.767
                    Skew: -0.457
                                         Prob(JB): 0.0339
                  Kurtosis: 3.584
                                        Cond. No.
                                                    5.23
```

Notes:

- [1] R² is computed without centering (uncentered) since the model does not contain a constant.
- [2] Standard Errors assume that the covariance matrix of the errors is correctly specified.

14. Build a model to study the impact of malnutrition and diseases on life expectancy. Identify the insignificant variables from the model. Use 'alpha = 1' if required.

```
In []: # type your code here
X = df.iloc[:,[3,8,12,11]]
X = sm.add_constant(X)
y = df['Life_Expectancy']

MLR_model4 = sm.OLS(y_train,X_train).fit()
```

```
In [117]: # type your code here
          # type your code here
          X = df.iloc[:,[3,8,12,11]]
          # set the dependent variable
          y = df['Life_Expectancy']
          # standardize all the columns of df_lifeExp
          X = X_scaler.fit_transform(X)
          # standardize the target explicitly and storing it in a new variable 'y'
          y = (y - y.mean()) / y.std()
          # split data into train subset and test subset
          # set 'random state' to generate the same dataset each time you run the code
          # 'test_size' returns the proportion of data to be included in the test set
          X_train, X_test, y_train, y_test = train_test_split(X, y, random_state = 30, test_size = 0.2)
In [118]: | lasso = Lasso(alpha = 1)
          lasso_model = lasso.fit(X_train, y_train)
In [119]: |MLR = LinearRegression()
          lin_model = MLR.fit(X_train, y_train)
In [120]: # Train
          train_pred = lin_model.predict(X_train)
          mse_train = mean_squared_error(y_train,train_pred)
          rmse_train = round(np.sqrt(mse_train),4)
          print(rmse_train)
          # Test
          test_pred = lin_model.predict(X_test)
          mse_test = mean_squared_error(y_test, test_pred)
          rmse_test = round(np.sqrt(mse_test),4)
          print(rmse_test)
          0.6853
          0.7967
In [121]: # Train
          train_pred = lasso_model.predict(X_train)
          mse_train = mean_squared_error(y_train,train_pred)
          rmse_train = round(np.sqrt(mse_train),4)
          print(rmse_train)
          # Test
          test_pred = lasso_model.predict(X_test)
          mse_test = mean_squared_error(y_test,test_pred)
          rmse_test = round(np.sqrt(mse_test),4)
          print(rmse_test)
          1.0039
          0.9813
                  15. Build a full model using the elastic net regression. Use 'alpha = 0.1' and 'l1 ratio=0.001'. Also, compute the RMSE for
                  train and test sets.
```

```
In [123]: # type your code here
# type your code here
# type your code here
X = df.iloc[:,:15]

# set the dependent variable
y = df['Life_Expectancy']

# standardize all the columns of df_LifeExp
X = X_scaler.fit_transform(X)

# standardize the target explicitly and storing it in a new variable 'y'
y = (y - y.mean()) / y.std()

# split data into train subset and test subset
# set 'random_state' to generate the same dataset each time you run the code
# 'test_size' returns the proportion of data to be included in the test set
X_train, X_test, y_train, y_test = train_test_split(X, y, random_state = 110, test_size = 0.2)
```

```
In [124]: | e_net = ElasticNet(alpha=0.1,l1_ratio=0.001)
          e_net_model = e_net.fit(X_train, y_train)
In [125]: # Train
          train_pred = e_net_model.predict(X_train)
          mse_train = mean_squared_error(y_train,train_pred)
          rmse_train = round(np.sqrt(mse_train),4)
          print(rmse_train)
          # Test
          test_pred = e_net_model.predict(X_test)
          mse_test = mean_squared_error(y_test,test_pred)
          rmse_test = round(np.sqrt(mse_test),4)
          print(rmse_test)
          0.3581
          0.3544
                  16. Build a full model. Identify the features that are significantly influencing the life expectancy. Use 'alpha = 0.5' if
           required.
In [126]: # type your code here
          # type your code here
          # type your code here
          X = df.iloc[:,:15]
          # set the dependent variable
          y = df['Life_Expectancy']
          # standardize all the columns of df_lifeExp
          X = X_scaler.fit_transform(X)
          # standardize the target explicitly and storing it in a new variable 'y'
          y = (y - y.mean()) / y.std()
          # split data into train subset and test subset
          # set 'random_state' to generate the same dataset each time you run the code
          # 'test_size' returns the proportion of data to be included in the test set
          X_train, X_test, y_train, y_test = train_test_split(X, y, random_state = 110, test_size = 0.2)
In [127]: |MLR1 = LinearRegression()
          lin_model1 = MLR1.fit(X_train, y_train)
In [128]: # Train
          train_pred = lin_model1.predict(X_train)
          mse_train = mean_squared_error(y_train,train_pred)
          rmse_train = round(np.sqrt(mse_train),4)
          print(rmse_train)
          # Test
          test_pred = lin_model1.predict(X_test)
          mse_test = mean_squared_error(y_test,test_pred)
          rmse_test = round(np.sqrt(mse_test),4)
          print(rmse_test)
```

5. Grid Search

0.350.3606

17. If ridge regression is used to build a full model. Which value is the best alpha from the list of values?

Consider the alpha values: [7,8,9,10,11]

```
X = df.iloc[:,:15]
          # set the dependent variable
          y = df['Life_Expectancy']
          # standardize all the columns of df_lifeExp
          X = X_scaler.fit_transform(X)
          # standardize the target explicitly and storing it in a new variable 'y'
          y = (y - y.mean()) / y.std()
          # split data into train subset and test subset
          # set 'random_state' to generate the same dataset each time you run the code
          # 'test_size' returns the proportion of data to be included in the test set
          X_train, X_test, y_train, y_test = train_test_split(X, y, random_state = 110, test_size = 0.2)
In [132]: | tuned_param = [{'alpha':[7,8,9,10,11]}]
          ridge = Ridge()
          ridge_grid = GridSearchCV(estimator=ridge,param_grid=tuned_param,cv=5)
          ridge_grid.fit(X_train,y_train)
          print(ridge_grid.best_params_)
          {'alpha': 11}
  In [ ]:
                  18. If lasso regression is used to build a full model. Which value is the best alpha from the list of values?
                  Consider the alpha values: [7,8,9,10,11]
In [133]: # type your code here
          X = df.iloc[:,:15]
          # set the dependent variable
          y = df['Life_Expectancy']
          # standardize all the columns of df_lifeExp
          X = X_scaler.fit_transform(X)
          # standardize the target explicitly and storing it in a new variable 'y'
          y = (y - y.mean()) / y.std()
          # split data into train subset and test subset
          # set 'random_state' to generate the same dataset each time you run the code
          # 'test_size' returns the proportion of data to be included in the test set
          X_train, X_test, y_train, y_test = train_test_split(X, y, random_state = 110, test_size = 0.2)
In [134]: | tuned_param = [{'alpha':[7,8,9,10,11]}]
          lasso = Lasso()
          lasso_grid = GridSearchCV(estimator=lasso,param_grid=tuned_param,cv=5)
          lasso_grid.fit(X_train,y_train)
          print(lasso_grid.best_params_)
          {'alpha': 7}
                   19. Perform regression analysis to study the impact of malnutrition and diseases on life expectancy. Determine the
                   optimal value of alpha and mixing parameter if elastic net regression is used to build the model. Consider the following
                  list of values:
                   alpha: [0.1,0.5,1,1.5,2,2.5,3,3.5,4]
                  I1_ratio: [0.5,0.55,0.6,0.65,0.7,0.75,0.8,0.85,0.9]
In [136]: # type your code here
          X = df.iloc[:,[3,8,11,12]]
          y = df['Life_Expectancy']
          X_train, X_test, y_train, y_test = train_test_split(X, y, random_state = 10, test_size = 0.2)
```

In [130]: # type your code here

```
In [137]: | tuned_param = [{'alpha':[0.1,0.5,1,1.5,2,2.5,3,3.5,4],'l1_ratio':[0.5,0.55,0.6,0.65,0.7,0.75,0.8,0.85,0.9]}]
          e_net1 = ElasticNet()
          enet_grid = GridSearchCV(estimator=e_net1,param_grid=tuned_param,cv=5)
          enet_grid.fit(X_train,y_train)
          print(enet_grid.best_params_)
          {'alpha': 0.1, 'l1_ratio': 0.9}
                 20. Build a full model. Select the optimal value for elastic net mixing parameter if the alpha value is 0.8
           Consider the values: I1_ratio: [0.001, 0.005, 0.01, 0.05, 0.1, 0.5, 1]
In [142]: # type your code here
          X = df.iloc[:,:15]
          y = df['Life_Expectancy']
          X_train, X_test, y_train, y_test = train_test_split(X, y, random_state = 110, test_size = 0.2)
In [143]: | tuned_param = [{'ll_ratio':[0.001, 0.005, 0.01, 0.05, 0.1, 0.5, 1]}]
          e_net2 = ElasticNet(alpha=0.8)
          enet_grid2 = GridSearchCV(estimator=e_net2,param_grid=tuned_param,cv=5)
          enet_grid2.fit(X_train,y_train)
          print(enet_grid2.best_params_)
```

{'l1_ratio': 1}